

EventSourcing and CQRS

A DYNAMIC DUO THAT WILL SOLVE GLOBAL
WARMING

(OR MAYBE SOMETHING A BIT MORE REALISTIC)

neo

Gold Sponsors



SOLVERA

Part of **Accenture**



online
business systems



/LOTLINX/®



Community Supporters



RICHARDSON

Introductions

Dave White AKT, AKC, PSM, CSM, MVP Alumni
Principal Consultant – Depth Consulting

- In technology for over 20 years
- Cloud-native, DevOps, Architect, Developer
- Kanban/Agile Evangelist, Speaker
- WesternDevs member
- Alumni MS MVP (2014, 2015, 2016)



<http://linkd.in/giMxuw>



@AgileRamblings



<https://www.westerndevs.com>



<https://github.com/agileramblings>



dave@depthconsulting.ca

Agenda

- Why are we here?
- EventSourcing
- CQRS
- Putting it all together
- Final Thoughts



Why are we here?

THE APPLICATION STORY WE'VE ALL
(PROBABLY) HEARD BEFORE



Learn. Share. Laugh. Cry.

There once was an application...

- Started in about 2012
- ASP.NET MVC – .cshtml with JSON over HTTP RPC added along the way
- Single database
- Entity Framework
- 100 users or so
 - No meaningful customer integrations
 - Minimal reporting



7-8 years later...

- Still ASP.NET MVC – .cshtml with JSON over HTTP RPC added along the way
 - Many Telerik/KendoUI grids/components
- Single database
- Entity Framework still
- 3000 users that include people, devices and load from third-party vendor integrations
 - > \$300 million dollars in annual revenue flow through this application
- BI ETLs run at night
- Large B2B EDI transactions run twice daily from major (read: \$\$\$) customers



Sounds ok until...

- Users complain about application slowness *all the time*.
 - Some grids for LOB activities take minutes to load
 - Analysis suggests the database is the reason...
- Let's scale up the database hardware!
- Let's allow queries and other transactions to use uncommitted, in-progress transactional writes

```
SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED;
```

- Let's spin up a read-only replica to try to take some of the load
- Call in the consultants to optimize data access

It didn't work...

- Users continue to complain about application slowness *all the time*.
 - LOB activities still take minutes to load
 - Long-running nightly activities don't finish by morning business hours
 - Analysis suggests the database is still the problem...
- The databases cost between \$4k-6k per month in Azure
 - Always pressing up against the biggest memory SKU vms available in Azure
- Allowing transactions to interact creates data integrity problems
- No auditing. With 3000 users/agents and data integrity problems, this becomes a requirement. This may add some load to the database.
- Read-only replica isn't having the expected impact
- There is nothing the consultants can do with the current architecture and EF



The problem wasn't SQL Server.


We had to do something...

- Begin cleaving chunks out of the monolith
 - Don't build another monolith
 - Don't use the current monoliths database either
- Within a microservice, re-evaluate how to do data access
 - Ensure flexibility with regard to distributing workloads
 - Auditability is a requirement
 - This isn't a talk about microservices
- Diverse consumers of business data need to be supported
 - REACT applications
 - ETL workloads
 - B2B service integrations
- Modern consumers have modern performance expectations



EventSourcing

DEMYSTIFYING AND MAKING IT LESS
SCARY



Choosing EventSourcing as a
persistence mechanism will have a
profound impact on the overall
implementation of your application.

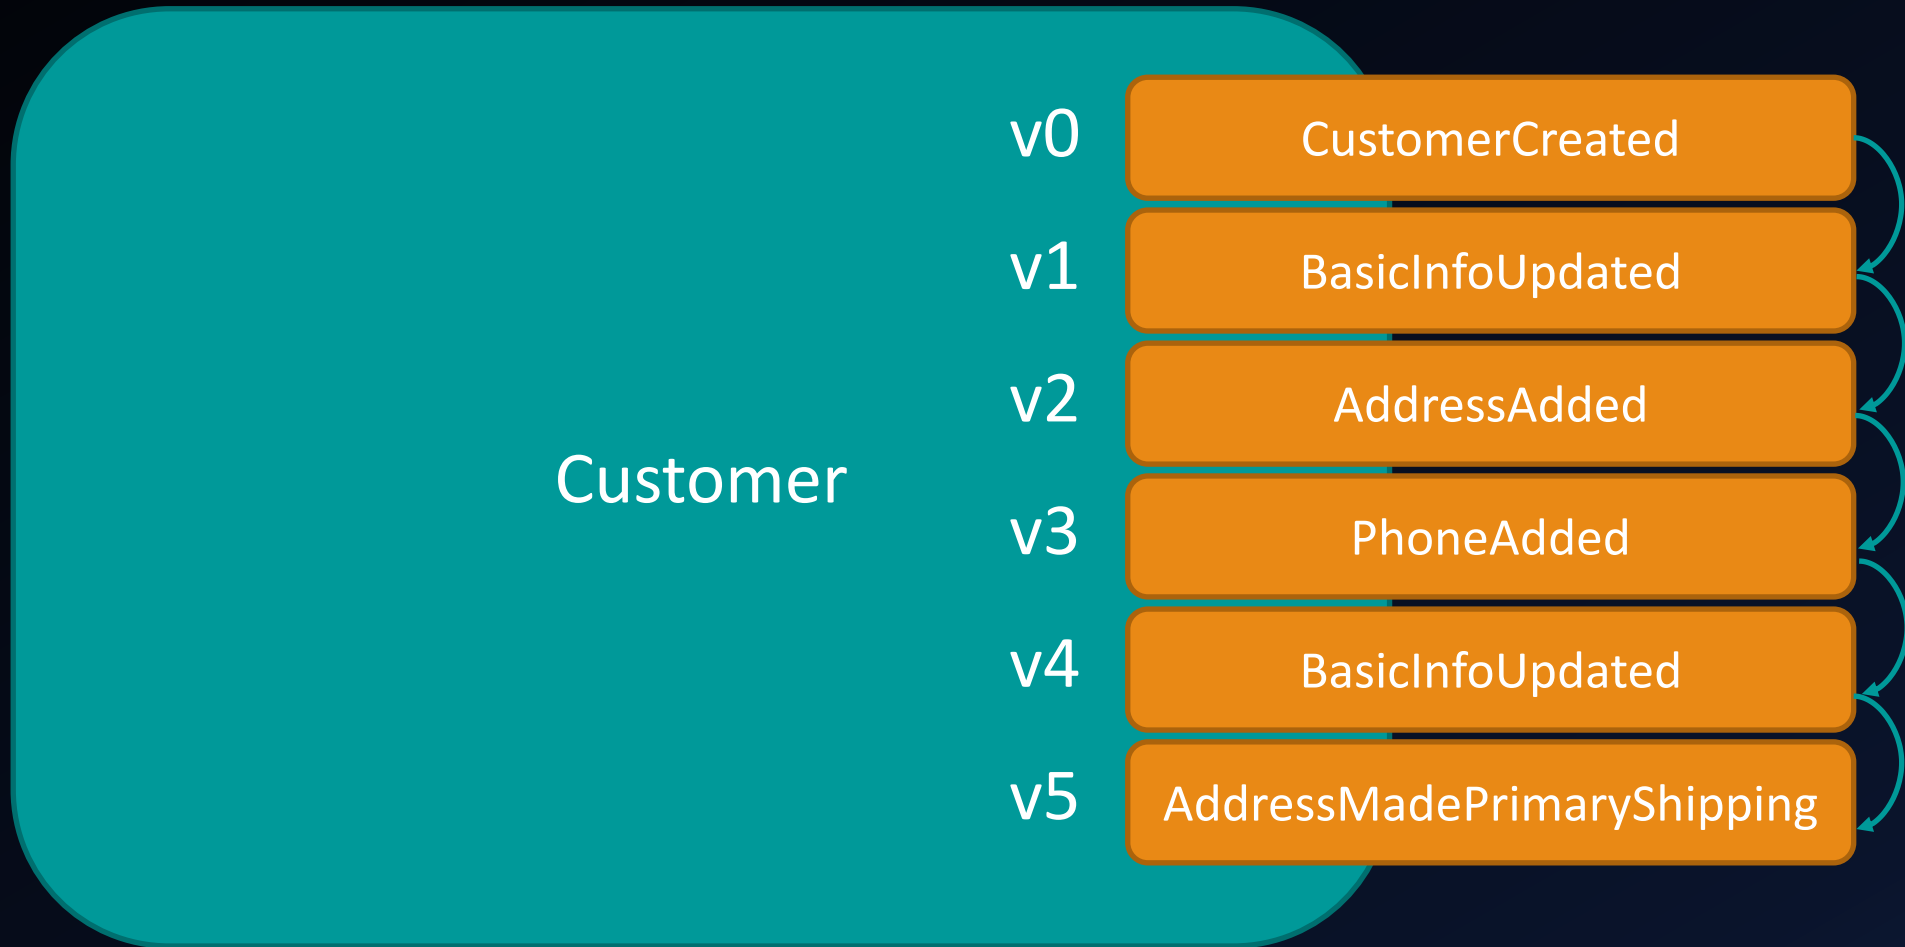
EventSourcing

- ... is a different style for saving System of Record business data
- ... aligns with alternate object modelling techniques (DDD)
- ... provides an intuitive approach to modelling important business objects in code
- ... is a bit more work
- ... may support auditing requirements by default
- ... won't bite you

Technically speaking...

- Stream of events describing what happened to an object
- The stream is (should be) append-only. Events are immutable.
- Each event represents a new “version” of the object
 - Events happen at a point in time
- Re-hydrating the object means reading the stream and applying events one at a time in order

A Visual



Event Store Records

SQLQuery2.sql - 127...ApiEvents (sa (68))* SQLQuery1.sql - not connected

<

vs Relational Records

SQLQuery4.sql - 127...piStaging (sa (70))* X SQLQuery2.sql - 127...ApiEvents (sa (68))*

/****** Script for SelectTopNRows command from SSMS *****/

```
SELECT TOP (1) [AggregateId]
, [Version]
, [FirstName]
, [LastName]
, [MiddleName]
, [DateOfBirth]
, [Email]
, [Gender]
, [CreatedBy]
, [CreatedOn]
, [ModifiedBy]
, [LastModifiedOn]
FROM [CustomerApiStaging].[dbo].[CustomerDetails]
WHERE AggregateId = '0a362995-7d3b-4233-9656-94aaf6ff92db'
```

100 %

Results Messages

	AggregateId	Version	FirstName	LastName	MiddleName	DateOfBirth
1	0a362995-7d3b-4233-9656-94aaf6ff92db	29	Marissa	Pearson	cosby	1979-10

SQLQuery5.sql - 127...piStaging (sa (74))* X SQLQuery4.sql - 127...piStaging (sa (70))* SQLQuery2.sql - 127...ApiEvents (sa (68))*

/****** Script for SelectTopNRows command from SSMS *****/

```
SELECT TOP (1000) [PhoneId]
, [IsPrimary]
, [Number]
, [CountryCode]
, [PhoneType]
, [ParentId]
FROM [CustomerApiStaging].[dbo].[Phones]
WHERE ParentId = '0a362995-7d3b-4233-9656-94aaf6ff92db'
```

100 %

Results Messages

	PhoneId	IsPrimary	Number	CountryCode	PhoneType	ParentId
1	D3F18523-5DD0-42D0-BC73-2606EA920EC1	0	447205029	4	occupy	0a362995-7d3b-4233-9656-94aaf6ff92db
2	A1188DEE-4166-4D57-A819-2ADB020F4475	0	239902877	6	Landlines	0a362995-7d3b-4233-9656-94aaf6ff92db
3	E5356942-2					
4	740BB68A-7					
5	D000979E-5					
6	30152DF8-3					
7	05F7947E-E					
8	A3AA9784-A					
9	DC3C787A-I					
10	70B91F10-1					
11	B93D489B-C					
12	0D5D7F9A-8					
13	DC549929-6					
14	6DC7DB65-7					
15	9FA23152-2					
16	334E6846-4					
17	16469B67-1					
18	650E71F4-F					
19	C47DF7FB-I					
20	3BD6E292-C					
21	D1D27E9A-C					
22	054860DF-9					

SQLQuery6.sql - 127...piStaging (sa (73))* X SQLQuery5.sql - 127...piStaging (sa (74))* SQLQuery4.sql - 127...piStaging (sa (70))*

/****** Script for SelectTopNRows command from SSMS *****/

```
SELECT TOP (1000) [AddressId]
, [IsPrimaryShipping]
, [IsPrimaryBilling]
, [Address1]
, [Address2]
, [City]
, [Region]
, [Country]
, [PostalCode]
, [ParentId]
FROM [CustomerApiStaging].[dbo].[Addresses]
WHERE ParentId = '0a362995-7d3b-4233-9656-94aaf6ff92db'
```

100 %

Results Messages

	AddressId	IsPrimaryShipping	IsPrimaryBilling	Address1	Address2	City	Region	Country	PostalCode	ParentId
1	331B8369-B39D-4B71-96F9-A8E5A50D11B4	1	0	8770 Paerdegat 12th Street						0a362995-7d3b-4233-9656-94aaf6ff92db
2	0EA8C149-B193-4333-85A7-CA877F1B8C8E	0	0	7814 Pine Street						0a362995-7d3b-4233-9656-94aaf6ff92db
3	1E2080B3-900A-4FCE-ABE9-EC02A0EB4E24	0	1	5627 Pembroke Street						0a362995-7d3b-4233-9656-94aaf6ff92db

One Schema to Rule them All

- Adding data to an event stream is as simple as adding events types, stored in the stream
- Adding data to a relational model usually means new tables or discriminators (yuck!) to an ever widening table

A brief word about consistency...

- Immediately consistent (state) vs. eventually consistent
- Objects in memory are the only place where we can be certain of being **immediately** consistent
- Everything else should be considered **eventually** consistent
- Anywhere that data is at rest is **eventually** consistent
 - First one to save wins!
- This is why versioning of domain objects is valuable

(Domain) Event handling too!

- With event-sourced persistence, your application can respond to specific events when persisting new items to the stream
- Handlers are decoupled from workflows in code
 - Register event handlers and they are invoked by an in-process message bus
 - Projectors and Reactors are two broad categories of event handlers
 - Projectors emit projections (point in time picture of a business object)
 - Reactors initiate an out-of-process interaction (Post to API, place a message on a external message bus)

Pros/Cons

PROS

- Intuitive object modeling
- What to test (scenarios) are easier to think about
- Easy (read: none) schema management
- Built-in auditing/history retention
- Version control of business objects
- Supports more successful cloud-scale (distributed) workloads
- Write-heavy workloads required extra care

CONS

- More data (de-normalized) usually means more storage costs
- Acquisition costs (queries) from event stream are may be more expensive
 - Mitigated with snapshots/mementos
- Cannot easily query data directly anymore
- Objects need to support old events
- More complexity
- Different/novel



CQRS

SPREADING THE LOAD OVER SPACE AND
TIME. AND... 42.





CQRS is not an architecture. It is a mindset.

CQRS

- ... Command Query **Responsibility** Segregation
- ... does *not* have to go hand in hand with EventSourcing
 - But they are a powerful combination
 - Always there are two...
- ... provides an intuitive approach that guides fit-for-purpose persistence implementations to emerge
- ... is a bit more work
- ... won't bite you

Commands

- Commands are a request to change the state of the world
- This usually means some sort of “write” activity
- e.g.
 - Please create a Customer record
 - Please add a phone number to a customer record

Query

- Queries are a request to see the state of the world
- This should not have any side-effects that change the state of the world
 - These side-effects are commands in disguise
- Queries (and responses) can (should?) be designed for specific consumers
 - JSON consumers
 - ETL integrations
 - Reporting tools
- e.g.
 - Get a customer record
 - Get a summary record of customer address/phone usage

Why should we keep these segregated?

- Ratio of workloads informs architectural design choices
 - Write workloads can be optimized/tuned
 - Query workloads can be optimized/tuned
 - May be deployed separately
- Can mitigate the “single persistence mechanism instance” risk
- Encourages fit-for-purpose query stores
 - NoSQL/Document databases
 - TimeSeries databases
 - Relational databases that are not attached the event store
 - Graph databases

Pros/Cons

PROS

- Encourages fit-for-purpose persistence mechanisms
- Separating command specific loads from query load allows provisioning “the right amount” of persistence compute
- Encourages architecturally consistent decisions
- Write-heavy workloads required extra care
- Query-heavy workloads can really benefit

CONS

- More data (de-normalized) usually means more storage costs
- More complexity
- Different/novel



Putting it all together

DID IT WORK THE WAY WE EXPECTED?



I have code!!!



Final Thoughts

SOME OF MINE, SOME OF YOURS

My thoughts

- I am a proponent of EventSourcing and CQRS and I am hoping that you will be too ... someday.
- Someday will come if you aren't intimidated by either of these techniques, which I am hoping you've taken away from today.
- It isn't an all or nothing adoption
- Neither ES or CQRS will bite you
 - any more than any other code you may write
- Any novel technology comes with sharp edges
 - Practice. Be Safe. Learn. Collaborate.

Questions

- If ES or CQRS is still scary, why?
- Technical questions?
- Usage concerns?
- Adoption concerns?

Resources

- <https://github.com/agileramblings>
 - PRDC2022 – arriving soon
 - Battleship-CQRS
- Eric Evans – Domain Driven Design – [Amazon.ca](https://www.amazon.ca)
- Vernon Vaughn – Implementing Domain Driven Design – [Amazon.ca](https://www.amazon.ca)
- Greg Young
 - https://cqrs.files.wordpress.com/2010/11/cqrs_documents.pdf
 - <https://www.youtube.com/watch?v=JHGkaShoyNs>
 - <https://www.eventstore.com/blog/author/greg-young>



Thank you